

Value Functions

Aaron Lou

June 2018

1 Omitting Policy Gradient

If we omit the policy gradient, then we can just consider our $A^\pi(s_t, a_t)$ and take the optimal policy in particular, our policy would be

$$\pi'(a_t | s_t) = \begin{cases} 1 & \text{if } a_t = \arg \max_{a_t} A^\pi(s_t, a_t) \\ 0 & \text{otherwise} \end{cases}$$

and we fit A^π in this way.

2 Policy and Value Iteration

2.1 Policy Iteration

The above method is called Policy Iteration. It works by

1. Evaluate A^π
2. Set $\pi \leftarrow \pi'$

and we calculate A^π as in actor-critic algorithms with

$$A^\pi(s, a) = r(s, a) + \gamma E(|V^\pi(s')| - V^\pi(s))$$

2.2 Dynamic Programming

By assuming our set of states and actions is relatively small (and discrete), then we can enumerate our V^π in a table. We use bootstrap update to get

$$V^\pi(s) \leftarrow E_{a \sim \pi(a|s)}[r(s, a) + \gamma E_{s' \sim p(s'|s, a)}[V^\pi(s')]]$$

and use this to evaluate $V^\pi(s)$. We can also just calculate the raw values of Q^π since it's the same thing as A^π . Our new value iteration method is

1. Set $Q(s, a) \leftarrow r(s, a) + \gamma E_{s' \sim p(s'|s, a)}[V(s')]$.
2. Set $V(s) \leftarrow \max_a Q(s, a)$

2.3 Fitted Value Iteration

We create a neural network to represent $V : S \rightarrow \mathbb{R}$ with parameters ϕ . We represent our loss as

$$\mathcal{L}(\phi) = \frac{1}{2} \|V_\phi(s) - \max_a Q^\pi(s, a)\|^2$$

We fit our values and get the new fitted value iteration method

1. Set $y_i \leftarrow \max_{a_i} (r(s_i, a_i) + \gamma E[V_\phi(s'_i)])$.
2. Set $\phi \leftarrow \arg \min_\phi \frac{1}{2} \sum_i \|V_\phi(s_i) - y_i\|^2$.

3 Q methods

3.1 Q Iteration

We can fit our Q value, which, as it turns out, doesn't require knowledge of transition dynamics

$$Q^\pi(s, a) \leftarrow r(s, a) + \gamma E_{s' \sim p(s'|s, \pi(s))} [Q^\pi(s', \pi(s'))]$$

And our policy iteration updates to

1. Evaluate $Q^\pi(s, a)$
2. Set $\pi \leftarrow \pi'$

3.2 Fitted Q Iteration

Our fitted form (using a neural net) iteration algorithm is as follows

1. Set $y_i \leftarrow r(s_i, a_i) + \gamma E[V_\phi(s'_i)]$
2. Set $\phi \leftarrow \arg \min_\phi \frac{1}{2} \sum_i \|Q_\phi(s_i, a_i) - y_i\|^2$

Where we can approximate our value of $E[V_\phi(s'_i)]$ with $\max_{a'} Q_\phi(s'_i, a'_i)$. Although this is a good off-policy method, it doesn't theoretically converge. The full algorithm is listed below

1. Collect dataset $\{(s_i, a_i, s'_i, r_i)\}$ using some policy
2. Set $y_i \leftarrow r(s_i, a_i) + \gamma \max_{a'_i} Q_\phi(s'_i, a'_i)$
3. Set $\phi \leftarrow \arg \max_\phi \frac{1}{2} \sum_i \|Q_\phi(s_i, a_i) - y_i\|^2$

where our parameters are, respectively,

1. N , our data set size.
2. K the number of iterations (we repeat steps 2 and 3 K times).
3. S the number of gradient steps.

This method is off-policy as step 1 is looking for any values, and our iteration doesn't depend on our current policy.

Our value ε is our error, and we have

$$\varepsilon = \frac{1}{2} E_{(s,a) \sim \beta} [Q_\phi(s, a) - [r(s, a) + \gamma \max_{a'} Q_\phi(s', a')]]$$

If $\varepsilon = 0$ then $Q_\phi(s, a) = r(s, a) + \gamma \max_{a'} Q_\phi(s', a')$. This is an optimal Q function, which corresponds to an optimal policy π' .

3.3 Online Q-learning algorithms

We can update our Q iteration algorithm as follows

1. Take some action a_i and observe (s_i, a_i, s'_i, r_i) .
2. $y_i = r(s_i, a_i) + \gamma \max_{a'_i} Q_\phi(s'_i, a'_i)$
3. $\phi \leftarrow \phi - \alpha \frac{dQ_\phi}{d\phi}(s_i, a_i)(Q_\phi(s_i, a_i) - y_i)$

In practice, we don't want to have the characteristic function as our first step, as this would remove all other possibilities. Instead, we use

$$\pi(a_t | s_t) = \begin{cases} 1 - \epsilon & \text{"epsilon greedy"} \\ \frac{\epsilon}{|\mathcal{A}|-1} & \end{cases}$$

$$\pi(a_t | s_t) \propto \exp(Q_\phi(s_t, a_t)) \text{ "Boltzmann exploration"}$$

4 Value Function Learning Theory

We define an operation $\mathcal{B} : \mathcal{B}V = \max_a r_a + \gamma \mathcal{T}_a V$. This is our update on $V(s)$ in our value iteration method. We notice there exists a best value V^* such that $\mathcal{B}V^* = V^*$. In particular

$$V^*(s) = r(s, a) + \gamma E[V^*(s')]$$

We also know that \mathcal{B} is a contraction (can be proved), so

$$\|\mathcal{B}V - \mathcal{B}\bar{V}\|_\infty \leq \gamma \|V - \bar{V}\|_\infty \implies \|\mathcal{B}V - \mathcal{B}V^*\|_\infty \leq \gamma \|V - V^*\|_\infty$$

so $V \leftarrow \mathcal{B}V$ goes to V^* . Furthermore, the gap always gets smaller by γ .

However, in the non-tabular case, our function approximator doesn't work out. In fact, if we reexamine our algorithm, we have

1. Set $y_i \leftarrow \max_{a_i} (r(s_i, a_i) + \gamma E[V_\phi(s'_i)])$.
2. Set $\phi \leftarrow \arg \min_\phi \frac{1}{2} \sum_i \|V_\phi(s_i) - y_i\|^2$.

We notice $\Pi V = \arg \min_{V' \in \Omega} \frac{1}{2} \sum \|V(s) - V'(s)\|^2$. Notice that Π is a contraction of the l_2 norm. However, $\Pi \mathcal{B}$ is not a general contraction, so it doesn't necessarily have to go to an optimal value.

Similarly, for fitted Q-iteration, this isn't a general contraction so it won't converge to our V^* . Online Q-learning isn't gradient descent since the target value $[r(s_i, a_i) + \gamma \max_{a'} Q_\phi(s'_i, a'_i)]$ is constantly changing. As a corollary, batch actor-critic algorithm suffers from the same problem.